



Comparing AI Models Across CFP® Domains

Objective: To compare the performance of leading AI language models – *OpenAI's ChatGPT (GPT-4)*, *Anthropic Claude*, *Meta's LLaMA-3*, and *Google Gemini* – across the eight knowledge domains of the CFP® (Certified Financial Planner) curriculum, using a perfect human expert score of 100 as a baseline.

CFP® Exam Domains and Weightings

The CFP® Certification Exam covers eight principal knowledge domains, each representing a portion of the exam. The domains and their approximate weightings are as follows ¹ :

- **Professional Conduct and Regulation** — 8%
- **General Principles of Financial Planning** — 15%
- **Risk Management and Insurance Planning** — 11%
- **Investment Planning** — 17%
- **Tax Planning** — 14%
- **Retirement Savings and Income Planning** — 18%
- **Estate Planning** — 10%
- **Psychology of Financial Planning** — 7%

These percentages indicate the relative importance of each topic area on the exam. In our comparison, a *Perfect Score* model would earn full marks in each domain (e.g. 8 points in Professional Conduct, 15 in General Principles, etc.), totaling **100** points.

AI Models and Benchmark Performance

ChatGPT (GPT-4): OpenAI's GPT-4 model has demonstrated near-human-level proficiency on many professional and academic benchmarks. For instance, GPT-4 recently passed the CPA (accounting) exam with an **85.1%** average score, a dramatic improvement over GPT-3.5's 53.1% (which failed) ² . It also scored in the top 10% of test-takers on a simulated bar exam and achieves roughly **80–90%** accuracy on advanced science and reasoning questions in evaluations ³ . These results suggest GPT-4 can attain roughly **80–90% of a perfect score** on complex, domain-specific tasks – making it a strong proxy for our "AI" performance in the financial planning context.

Anthropic Claude: Claude (latest version in 2025, often referred to as Claude 3 or "Claude Sonnet") is another leading model, tuned for extended reasoning. Claude has shown **impressive performance** on reasoning and knowledge benchmarks, in some cases **rivaling or even exceeding GPT-4**. For example, Claude 3 scored about **86.8%** on a broad knowledge test (undergraduate level) – on par with GPT-4 – and significantly outperformed GPT-4 on a graduate-level reasoning benchmark (50.4% vs 35.7% accuracy) ⁴ . Claude 3 also excelled at coding tasks in one evaluation, with an **84.9%** success rate compared to GPT-4's 67.0% ⁵ . Overall, Claude's capabilities appear to be in the same general range as GPT-4 for many tasks (roughly 70–85% of expert level, depending on the task), with particular strengths in complex reasoning.

Meta LLaMA-3: Meta's LLaMA 3 is an advanced open-source LLM, available in various model sizes (the largest with hundreds of billions of parameters). On standard benchmarks, LLaMA-3 can approach the closed-source models in general knowledge – for instance, a 70B-parameter LLaMA-3 scored **~82%** on the MMLU academic knowledge benchmark ⁶. However, it tends to lag in specialized or computation-heavy areas; the same 70B model managed only **around 30%** accuracy on a math word-problem test (MATH), versus GPT-4's 75+% on that test ⁶. Even a much larger 400B-parameter LLaMA-3 improved to ~57.8% on MATH, still trailing top models ⁶. These results indicate that while LLaMA-3 has strong general knowledge, its overall performance on complex financial-planning problems likely reaches **around 70-80% of expert level** at best. (In other words, it's very capable for an open model, but still a notch below GPT-4/Claude in this domain.)

Google Gemini: Google's Gemini (first released in late 2023 and iteratively improved through 2025) is designed to compete directly with GPT-4. Early versions of Gemini did not immediately surpass GPT-4 on standard benchmarks – in one 2024 comparison, Gemini (Pro 1.5) scored ~81.9% on MMLU and ~58.5% on the MATH dataset, slightly behind GPT-4's performance ⁶. In fact, in that evaluation neither of the tested Gemini models took the top spot on any benchmark – GPT-4 and LLaMA-3 held those honors ⁷. However, Google has aggressively improved Gemini's capabilities. The largest configuration, **Gemini Ultra**, was reported to be the **first model to exceed human experts on MMLU** with a 90.0% score ⁸, and generally has achieved state-of-the-art results on many tasks. By 2025, a practical version of Gemini intended for broad use (Gemini 2.x) can be estimated to perform in the same ballpark as GPT-4. For our purposes, we'll assume Gemini's overall proficiency is roughly **80-85% of the perfect score** – comparable to GPT-4, with the expectation that it matches or slightly trails GPT-4 on financial reasoning tasks (pending more public benchmark data in finance).

Scoring Methodology

To compare these models, we define a **normalized 0-100 score scale**, where **100** represents a "Perfect Score" (essentially an expert human achieving 100% across all CFP domains). Each AI model is assigned a total score on this scale based on the available benchmark evidence for its financial and reasoning abilities.

- **ChatGPT (GPT-4):** ~85/100 (approximately 85% of a perfect expert performance, given its ~85% scores on finance-related exams ²).
- **Anthropic Claude:** ~80/100 (slightly below GPT-4, though it excels on some tasks, we estimate around 80% overall).
- **Google Gemini:** ~83/100 (estimated to be on par with GPT-4, perhaps just a bit behind in financial planning contexts, pending further data).
- **Meta LLaMA-3:** ~75/100 (the lowest of the group, as open models still trail the fine-tuned proprietary models in complex, domain-specific performance).

Distribution across domains: In the absence of detailed domain-by-domain performance data for each AI, we distribute each model's total score across the CFP domains *proportionally to the official weightings*. This assumes the model's strengths and weaknesses affect all domains evenly (which may not be true in reality, but proportional allocation provides a reasonable proxy given limited information). For example, ChatGPT's total of ~85 points is apportioned 8% to Professional Conduct (~6.8 points), 15% to General Principles (~12.8 points), and so on. We round these to whole numbers for clarity, so ChatGPT gets about 7 points in the first domain (out of 8 possible) and 13 in the second (out of 15), etc., summing to ~85. The table below summarizes the approximate **score breakdown by domain** for each model:

Estimated Scores by CFP® Domain (out of 100)

Model	Professional >(8%)	General >(15%)	Risk >(11%)	Investment >(17%)	Tax >(14%)	Retirement >(18%)	Estate >(10%)
ChatGPT (GPT-4)	7	13	9	14	12	15	8
Anthropic Claude	6	12	9	14	11	14	8
Meta LLaMA-3	6	11	8	13	10	14	8
Google Gemini	7	12	9	14	12	15	8
Perfect Score	8	15	11	17	14	18	10

(Points in each domain are weighted by that domain's percentage of the exam. "Perfect Score" corresponds to full marks in every category, totaling 100.)

Visualization

Figure: Stacked horizontal bar chart comparing AI models against a Perfect Score baseline across the 8 CFP® domains. Each bar is segmented according to the CFP domain weightings (color-coded to match the domain list above). The total length of each bar represents the model's overall score out of 100 (with "Perfect Score" reaching the full 100 points).

As shown in the chart, **ChatGPT (GPT-4)** comes the closest to the perfect 100-point benchmark, with roughly 85 points distributed across the domains, followed closely by **Google's Gemini** (~83). **Anthropic Claude** is not far behind at ~80, while **Meta's LLaMA-3** trails with around 75 points. In all cases, the models achieve strong performance in each domain relative to the maximum, but none yet reach the full 100 – highlighting that even the best AI still falls short of a true expert's level of knowledge and skill across all areas of financial planning. The proportional domain breakdown also illustrates that domains with higher weights (like Retirement Planning and Investment Planning) contribute more to each model's total score, mirroring the CFP exam's emphasis on those areas.

Sources: The CFP® domain weightings are from CFP Board/Kaplan exam outlines ¹. AI model performance estimates are based on reported benchmark results for GPT-4 ² ³, Claude ⁴ ⁵, LLaMA-3 ⁶, and Gemini ⁷ ⁸. Each model's score distribution is extrapolated from these benchmark percentages onto the CFP domain weights.

¹ How to Pass the CFP Exam | 5 Steps to Success | Kaplan

<https://www.kaplanfinancial.com/resources/getting-started/how-to-pass-the-cfp-exam>

2 ChatGPT aces CPA exam after prior version flunked | CFO Dive

<https://www.cfodive.com/news/chatgpt-aces-cpa-exam-prior-version-flunked-accounting-audit/651333/>

3 Comparative Business Study: Claude Models vs GPT-4

<https://www.fluid.ai/blog/comparative-business-study-claude-models-vs-gpt-4>

4 5 Claude 3 vs GPT 4: Is Claude better than GPT-4? | Merge

<https://merge.rocks/blog/claude-3-vs-gpt-4-is-claude-better-than-gpt-4>

6 7 The Battle of the LLMs: Meta's Llama 3 vs. GPT-4 vs. Gemini

<https://capestart.com/resources/blog/the-battle-of-the-llms-llama-3-vs-gpt-4-vs-gemini/>

8 Introducing Gemini: Google's most capable AI model yet

<https://blog.google/technology/ai/google-gemini-ai/>